

---

## **Lab 02: Map Reduce Programming**

CSC14118 Introduction to Big Data 20KHMT1

X-HAT

2023-03-31

# Contents

<b>1</b>	<b>Lab 02: Map Reduce Programming</b>	<b>3</b>
1.1	1. WordCount Program . . . . .	3
1.1.1	Step 1: Program's solution . . . . .	3
1.1.2	Step 2: Class Creation . . . . .	4
1.1.3	Step 3: Create directory structure for program in Hadoop . . . . .	5
1.1.4	Step 5: Final result . . . . .	5
1.2	2. WordSizeWordCount Program . . . . .	7
1.2.1	Step 1: Program's solution . . . . .	7
1.2.2	Step 2: Class Creation . . . . .	8
1.2.3	Step 3: Create directory structure for program in Hadoop . . . . .	8
1.2.4	Step 5: Final result . . . . .	9
1.3	3. WeatherData program . . . . .	11
1.3.1	Step 1: Program's solution . . . . .	11
1.3.2	Step 2: Class Creation . . . . .	12
1.3.3	Step 3: Create directory structure for program in Hadoop . . . . .	12
1.3.4	Step 5: Final result . . . . .	13
1.4	4. Patent Program . . . . .	13
1.4.1	Step 1: Program's solution . . . . .	13
1.4.2	Step 2: Class Creation . . . . .	16
1.4.3	Step 3: Create directory structure for program in Hadoop . . . . .	16
1.4.4	Step 4: Create Jar File and deploy it to Hadoop . . . . .	18
1.4.5	Step 5: Final result . . . . .	18
1.5	5. MaxTemp Program . . . . .	18
1.5.1	Step 1: Program's solution . . . . .	18
1.5.2	Step 2: Class Creation . . . . .	20
1.5.3	Step 3: Create directory structure for program in Hadoop . . . . .	21
1.5.4	Step 5: Final result . . . . .	21
1.6	6. AverageSalary Program . . . . .	22
1.6.1	Step 1: Program's solution . . . . .	22
1.6.2	Step 2: Class Creation . . . . .	24

1.6.3	Step 3: Create directory structure for program in Hadoop . . . . .	24
1.6.4	Step 4: Create Jar File and deploy it to Hadoop . . . . .	26
1.6.5	Step 5: Final result . . . . .	26
1.7	7. De Identify HealthCare Program . . . . .	26
1.7.1	Step 3: Program's solution . . . . .	26
1.7.2	Step 2: Class Creation . . . . .	29
1.7.3	Step 3: Create directory structure for program in Hadoop . . . . .	30
1.7.4	Step 4: Create Jar File and deploy it to Hadoop . . . . .	30
1.7.5	Step 5: Final result . . . . .	31
1.8	8 Music Track Program . . . . .	31
1.8.1	Step 1: Program's solution . . . . .	31
1.8.2	Step 2: Class Creation . . . . .	40
1.8.3	Step 3: Create directory structure for program in Hadoop . . . . .	40
1.8.4	Step 5: Final result . . . . .	41
1.9	9. Telecom Call Data Record Program . . . . .	48
1.9.1	Step 1: Program's solution . . . . .	48
1.9.2	Step 2: Class Creation . . . . .	50
1.9.3	Step 3: Create directory structure for program in Hadoop . . . . .	50
1.9.4	Step 5: Final result . . . . .	51
1.10	10. Count Connected Components Program . . . . .	51
1.10.1	Step 1: Program's solution . . . . .	51
1.10.2	Step 2: Class Creation . . . . .	54
1.10.3	Step 3: Create directory structure for program in Hadoop . . . . .	55
1.10.4	Step 4: Create Jar File and deploy it to Hadoop . . . . .	56
1.10.5	Step 5: Final result . . . . .	56
1.11	Self-reflection . . . . .	56
1.11.1	20127435 - Tran Van An . . . . .	56
1.11.2	20127395 - Phan Minh Xuan . . . . .	58
1.11.3	20127032 - Bui Gia Huy . . . . .	58
1.11.4	20127631 - Thai Van Thien . . . . .	58
1.12	References . . . . .	59

# 1 Lab 02: Map Reduce Programming

## 1.1 1. WordCount Program

### 1.1.1 Step 1: Program's solution

- Import:

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

- Mapper:

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

- Reducer:

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the StringTokenizer to separate the words in the value. Then it sends each word to the Reducer with a value of 1.
- In the reducer metho: This reducer will take the words from the Mapper and calculate the total number of occurrences of each word by adding the values of 1s together.

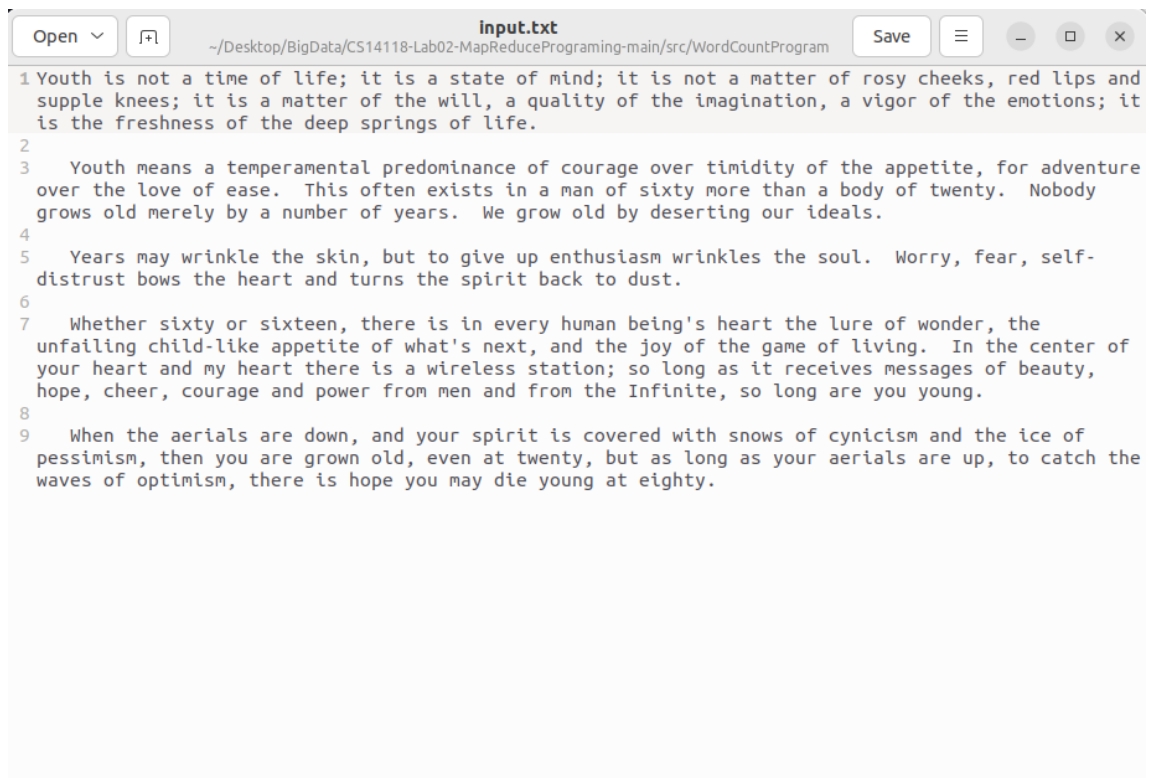
### 1.1.2 Step 2: Class Creation

```
jar -cvf WordCount.jar -C classes/ .
```

### 1.1.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /WordCount
hadoop fs -mkdir /WordCount/Input
hadoop fs -put 'local input file's path ' /WordCount/Input
```

- Example input:



- ### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" WordCount /WordCount/Input /WordCount/Output
```

### 1.1.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:
- 
-

/WordCount/Output Go!

Show  entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	pmx	supergroup	0 B	Mar 28 21:43	1	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	pmx	supergroup	1.16 KB	Mar 28 21:43	1	128 MB	part-r-00000	

Showing 1 to 2 of 2 entries

Previous **1** Next

Hadoop, 2022.

Figure 1.1: Output 1

Block information -- Block 0

Block ID: 1073741969  
Block Pool ID: BP-1926325839-127.0.1.1-1678203211420  
Generation Stamp: 1145  
Size: 1184  
Availability:

- pmx

File contents

```
In 1
Infinite, 1
Nobody 1
This 1
We 1
When 1
Whether 1
Worry, 1
```

Figure 1.2: Output 2

## 1.2 2. WordSizeWordCount Program

### 1.2.1 Step 1: Program's solution

- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {

            word.set(tokenizer.nextToken());
            String length = String.valueOf(word.getLength());
            Text len = new Text(length);
            context.write(len, one);
        }
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException

    {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
    }
}
```



```
    }  
    context.write(key, new IntWritable(sum));  
  }  
}
```

- Main:

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = new Job(conf, "WordSizeWordCount");  
    job.setJarByClass(WordSizeWordCount.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    job.setMapperClass(Map.class);  
    job.setCombinerClass(Reduce.class);  
    job.setReducerClass(Reduce.class);  
    job.setInputFormatClass(TextInputFormat.class);  
    job.setOutputFormatClass(TextOutputFormat.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    job.waitForCompletion(true);  
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the StringTokenizer to separate the words in the value. Then it sends each word's length to the Reducer with a value of 1.
- In the reducer method: This reducer will take the word's length from the Mapper and calculate the total number of occurrences of each word's length by adding the values of 1s together.

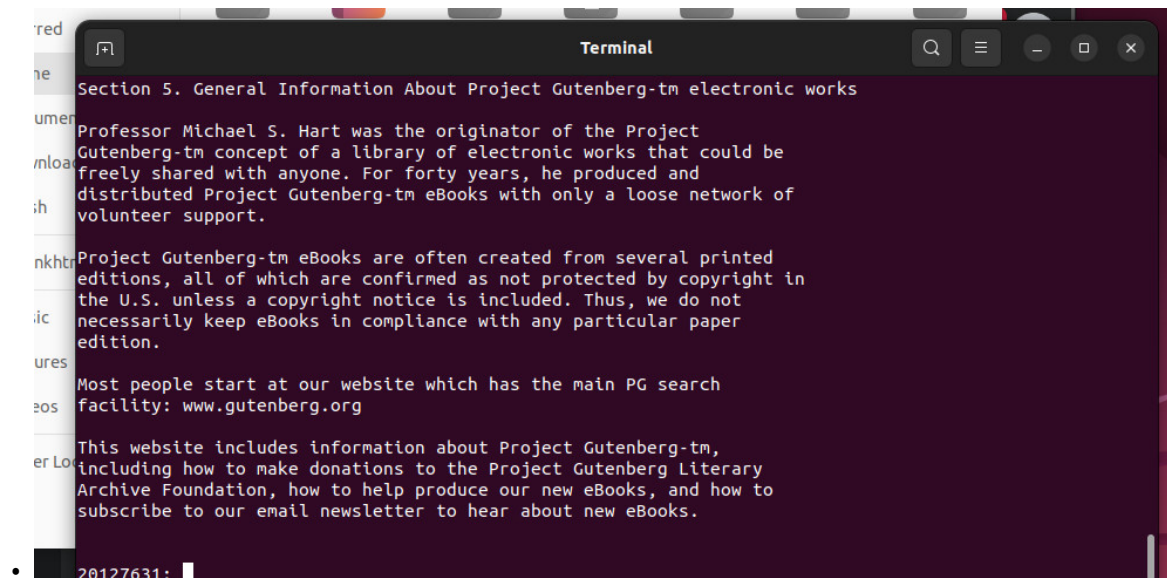
### 1.2.2 Step 2: Class Creation

```
jar -cvf WordSizeWordCount.jar -C classes/ .
```

### 1.2.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /WordSizeWordCount  
hadoop fs -mkdir /WordSizeWordCount  
hadoop fs -put 'local input file's path ' /WordSizeWordCount/Input
```

- Example input:



### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" WordSizeWordCount /WordSizeWordCount/Input
↪ /WordSizeWordCount/Output
```

### 1.2.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:

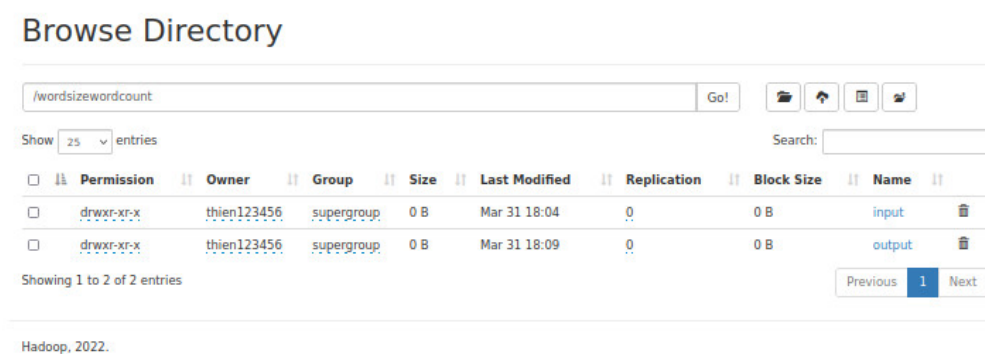
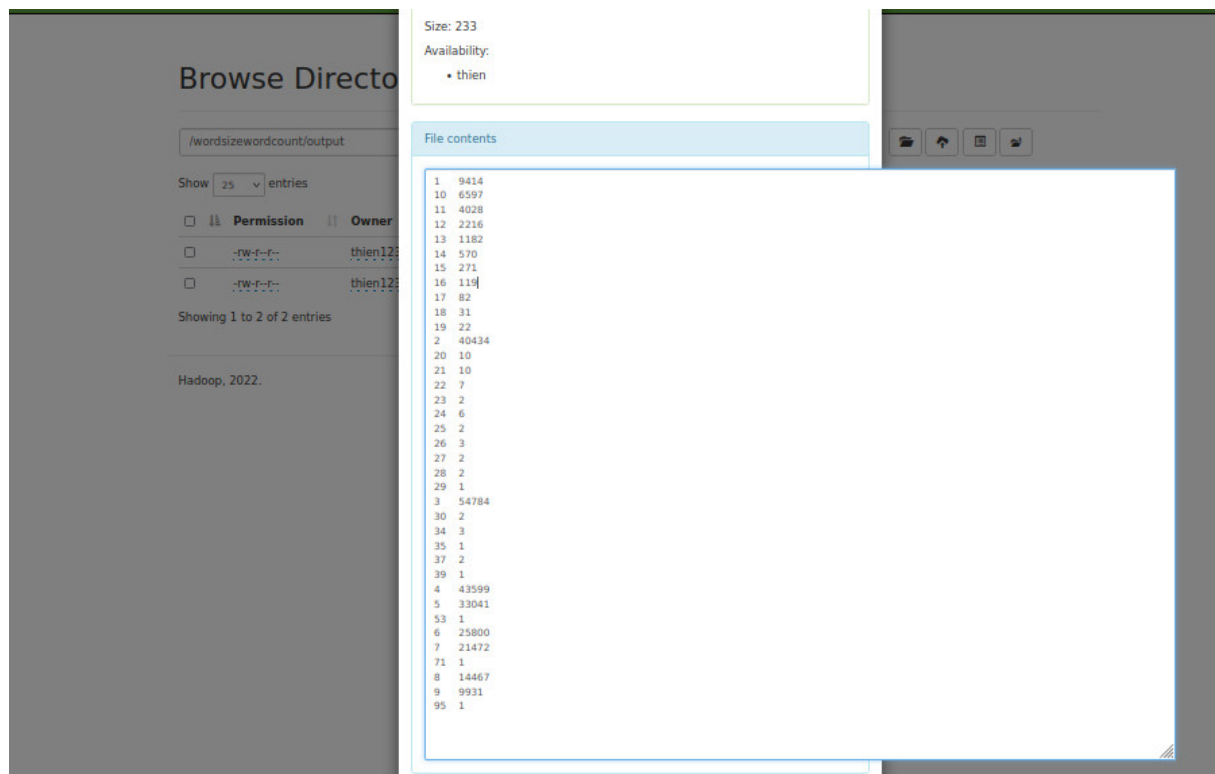


Figure 1.3: Output 1

- 
-

**Figure 1.4:** Output 2

## 1.3 3. WeatherData program

### 1.3.1 Step 1: Program's solution

- Import:

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

- Mapper:

```
public static class Map
    extends Mapper<Object, Text, Text, Text>{

    public void map(Object key, Text value, Context context
    ) throws IOException, InterruptedException {
        String line = value.toString();
        String[] tokens = line.split("\\s+");
        String date = tokens[1];
        float tempMax = Float.parseFloat(tokens[6].trim());
        float tempMin = Float.parseFloat(tokens[7].trim());

        if(tempMax > 40.0) {
            context.write(new Text("Hot Day " + date), new Text(String.valueOf(tempMax)));
        }
        if(tempMin < 10.0) {
            context.write(new Text("Cold Day " + date), new
↵ Text(String.valueOf(tempMin)));
        }
    }
}
```

- Reducer:

```
public static class Reduce
    extends Reducer<Text,Text,Text,Text>{

    public void reduce(Text key, Iterator<Text> values,
        Context context
    ) throws IOException, InterruptedException {
        String temperature = values.next().toString();
        context.write(key, new Text(temperature));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "weather");
    job.setJarByClass(WeatherData.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the `split()` method to separate the fields in a line (inside is a regular expression indicate one or more whitespace). It then extract the date, minTemp and maxTemp (at index 1, 6, 7 respectively). Then, it check for the conditions (`tempMax > 40` or `tempMin < 10`) and write the coressponding output to the context.
- In the reducer metho: This reducer will take the words from the Mapper and print out the result.

### 1.3.2 Step 2: Class Creation

```
jar -cvf WeatherData.jar -C classes/ .
```

### 1.3.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /Weather
hadoop fs -mkdir /Weather/Input
hadoop fs -put 'local input file's path ' /Weather/Input
```

- Example input:

weather\_data.txt - Notepad

23907	20150101	2.423	-98.08	30.62	2.2	-0.6	0.8	0.9	6.2	1.47	C	3.7	1.1	2.5	99.9	85.4	97.2	0.369	0.308	-99.000	-99.000	-99.000	7.0
8.1	-9999.0	-9999.0	-9999.0	30.62	3.5	1.3	2.4	2.2	9.0	1.43	C	4.9	2.3	3.1	100.0	98.8	99.8	0.391	0.327	-99.000	-99.000	-99.000	7.1
7.9	-9999.0	-9999.0	-9999.0	30.62	15.9	2.3	9.1	7.5	2.9	11.00	C	16.4	2.9	7.3	100.0	34.8	73.7	0.450	0.397	-99.000	-99.000	-99.000	7.6
7.9	-9999.0	-9999.0	-9999.0	30.62	9.2	-1.3	3.9	4.2	0.0	13.24	C	12.4	-0.5	4.9	82.0	40.6	61.7	0.414	0.352	-99.000	-99.000	-99.000	7.3
7.9	-9999.0	-9999.0	-9999.0	30.62	10.9	-3.7	3.6	2.6	0.0	13.37	C	14.7	-3.0	3.8	77.9	33.3	57.4	0.399	0.340	-99.000	-99.000	-99.000	6.3
7.0	-9999.0	-9999.0	-9999.0	30.62	20.2	2.9	11.6	10.9	0.0	12.90	C	22.0	1.6	9.9	67.7	30.2	49.3	0.395	0.335	-99.000	-99.000	-99.000	8.0
8.0	-9999.0	-9999.0	-9999.0	30.62	10.9	-3.4	3.8	4.5	0.0	12.68	C	12.4	-2.1	5.5	82.7	36.5	55.7	0.387	0.328	-99.000	-99.000	-99.000	7.6
8.3	-9999.0	-9999.0	-9999.0	30.62	0.6	-7.9	-3.6	-3.3	0.0	4.98	C	3.9	-4.8	-0.5	57.7	37.6	48.1	0.372	0.316	-99.000	-99.000	-99.000	4.7
6.1	-9999.0	-9999.0	-9999.0	30.62	2.0	0.1	1.0	0.8	0.0	2.52	C	4.1	1.2	2.5	87.8	48.9	64.4	0.368	0.312	-99.000	-99.000	-99.000	5.4
6.2	-9999.0	-9999.0	-9999.0	30.62	0.5	-2.0	-0.8	-0.6	3.3	2.11	C	2.5	-0.1	1.4	99.9	47.7	85.8	0.373	0.314	-99.000	-99.000	-99.000	5.1
6.0	-9999.0	-9999.0	-9999.0	30.62	10.9	0.0	5.4	4.4	2.9	6.38	C	12.7	1.3	5.8	100.0	77.8	97.1	0.420	0.362	-99.000	-99.000	-99.000	6.5
6.7	-9999.0	-9999.0	-9999.0	30.62	6.5	1.4	4.0	4.3	0.0	1.55	C	6.9	2.7	5.1	100.0	89.4	97.8	0.412	0.350	-99.000	-99.000	-99.000	7.3
7.5	-9999.0	-9999.0	-9999.0	30.62	3.0	-0.7	1.1	1.2	0.0	3.26	C	5.6	0.7	2.9	99.7	80.7	90.7	0.401	0.337	-99.000	-99.000	-99.000	6.1
6.8	-9999.0	-9999.0	-9999.0	30.62	2.9	0.9	1.9	1.8	0.0	1.88	C	4.7	2.0	3.1	99.6	90.8	97.9	0.395	0.331	-99.000	-99.000	-99.000	6.1
6.7	-9999.0	-9999.0	-9999.0	30.62	13.2	1.2	7.2	6.4	0.0	13.37	C	16.4	1.4	6.7	98.9	46.7	73.4	0.395	0.333	-99.000	-99.000	-99.000	6.7
7.0	-9999.0	-9999.0	-9999.0	30.62	16.7	3.5	10.1	9.9	0.0	13.68	C	19.2	1.3	8.7	80.2	38.1	58.2	0.391	0.330	-99.000	-99.000	-99.000	7.3
7.4	-9999.0	-9999.0	-9999.0	30.62	19.5	5.0	12.2	12.3	0.0	10.96	C	20.9	3.3	10.6	87.7	30.4	55.7	0.388	0.327	-99.000	-99.000	-99.000	8.7
8.4	-9999.0	-9999.0	-9999.0	30.62	20.9	7.6	14.3	13.7	0.0	15.03	C	23.4	3.5	11.9	45.9	14.6	31.4	0.383	0.325	-99.000	-99.000	-99.000	9.5
9.2	-9999.0	-9999.0	-9999.0	30.62	23.9	6.7	15.3	14.3	0.0	14.10	C	25.6	3.8	12.6	65.3	26.8	45.6	0.376	0.321	-99.000	-99.000	-99.000	9.9
9.5	-9999.0	-9999.0	-9999.0	30.62	26.0	9.5	17.8	15.9	0.0	14.57	C	27.9	6.5	14.5	88.4	16.1	50.2	0.373	0.320	-99.000	-99.000	-99.000	10.9
10.4	-9999.0	-9999.0	-9999.0	30.62	11.0	6.9	8.9	8.9	1.7	2.71	C	13.1	6.8	9.7	99.2	68.0	88.1	0.369	0.317	-99.000	-99.000	-99.000	10.7
10.6	-9999.0	-9999.0	-9999.0	30.62	8.6	3.5	6.1	5.6	39.6	1.28	C	9.1	4.1	6.3	99.6	95.2	98.0	0.546	0.418	-99.000	-99.000	-99.000	9.0
23907	20150122	2.423	-98.08	30.62	8.6	3.5	6.1	5.6	39.6	1.28	C	9.1	4.1	6.3	99.6	95.2	98.0	0.546	0.418	-99.000	-99.000	-99.000	9.0

Ln 1, Col 1 100% Unix (LF) UTF-8

### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" WeatherData /Weather/Input /Weather/Output
```

### 1.3.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:

- 
- 

## 1.4 4. Patent Program

### 1.4.1 Step 1: Program's solution

- Mapper:

## Browse Directory

/Weather/Output

Go!

Show

25

entries

Search:

<input type="checkbox"/>	<div><div></div></div> Permission	<div><div></div></div> Owner	<div><div></div></div> Group	<div><div></div></div> Size	<div><div></div></div> Last Modified	<div><div></div></div> Replication	<div><div></div></div> Block Size	<div><div></div></div> Name	<div><div></div></div>
<input type="checkbox"/>	<div>-rw-r--r--</div>	ADMIN	supergroup	0 B	Mar 31 20:53	1	128 MB	<a href="#">_SUCCESS</a>	<div><div></div></div>
<input type="checkbox"/>	<div>-rw-r--r--</div>	ADMIN	supergroup	864 B	Mar 31 20:53	1	128 MB	<a href="#">part-r-00000</a>	<div><div></div></div>

Showing 1 to 2 of 2 entries

Previous

1

Next

Hadoop, 2022.

**Figure 1.5:** Output 1

```
public static class PatentMapper
    extends Mapper<Object, Text, Text, Text> {
    Text k = new Text();
    Text v = new Text();

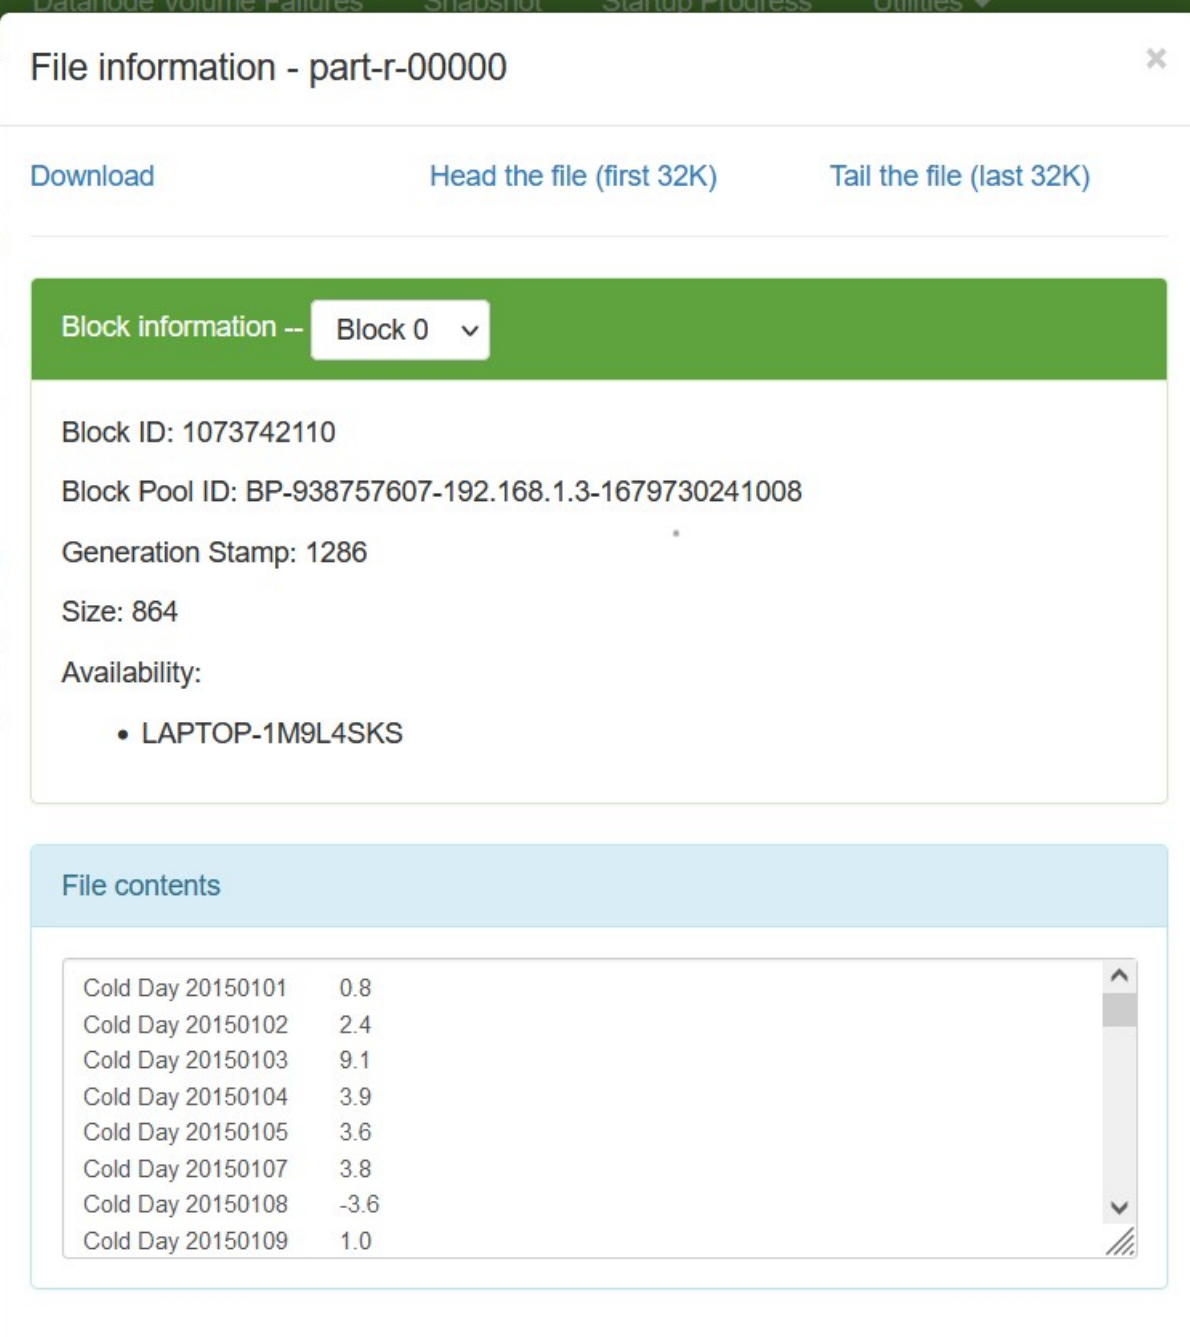
    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();

        StringTokenizer tokenizer = new StringTokenizer(line, " ");
        while (tokenizer.hasMoreTokens()) {
            String token = tokenizer.nextToken();
            k.set(token);
            String token1 = tokenizer.nextToken();
            v.set(token1);
            context.write(k, v);
        }
    }
}
```

- Reducer:

```
public static class SumSubPatentReducer
    extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (Text x : values) {
            sum++;
        }
    }
}
```



File information - part-r-00000

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0 ▾

Block ID: 1073742110

Block Pool ID: BP-938757607-192.168.1.3-1679730241008

Generation Stamp: 1286

Size: 864

Availability:

- LAPTOP-1M9L4SKS

File contents

Cold Day 20150101	0.8
Cold Day 20150102	2.4
Cold Day 20150103	9.1
Cold Day 20150104	3.9
Cold Day 20150105	3.6
Cold Day 20150107	3.8
Cold Day 20150108	-3.6
Cold Day 20150109	1.0

**Figure 1.6:** Output 2



```
        String result = Integer.toString(sum);
        context.write(key, new Text(result));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "patent program");
    job.setJarByClass(PatentProgram.class);

    job.setMapperClass(PatentMapper.class);
    job.setReducerClass(SumSubPatentReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- The main idea for this program is that collecting pair of token in Map function, after combining, we count them through the their key and write in output file.

### 1.4.2 Step 2: Class Creation

- After complete code in Java, we need to generate file jar from builded classes by below command:

```
jar -cvf PatentProgram.jar -C classes/ .
```

- Notice: Make sure that you export HADOOP\_CLASSPATH before buiding file jar

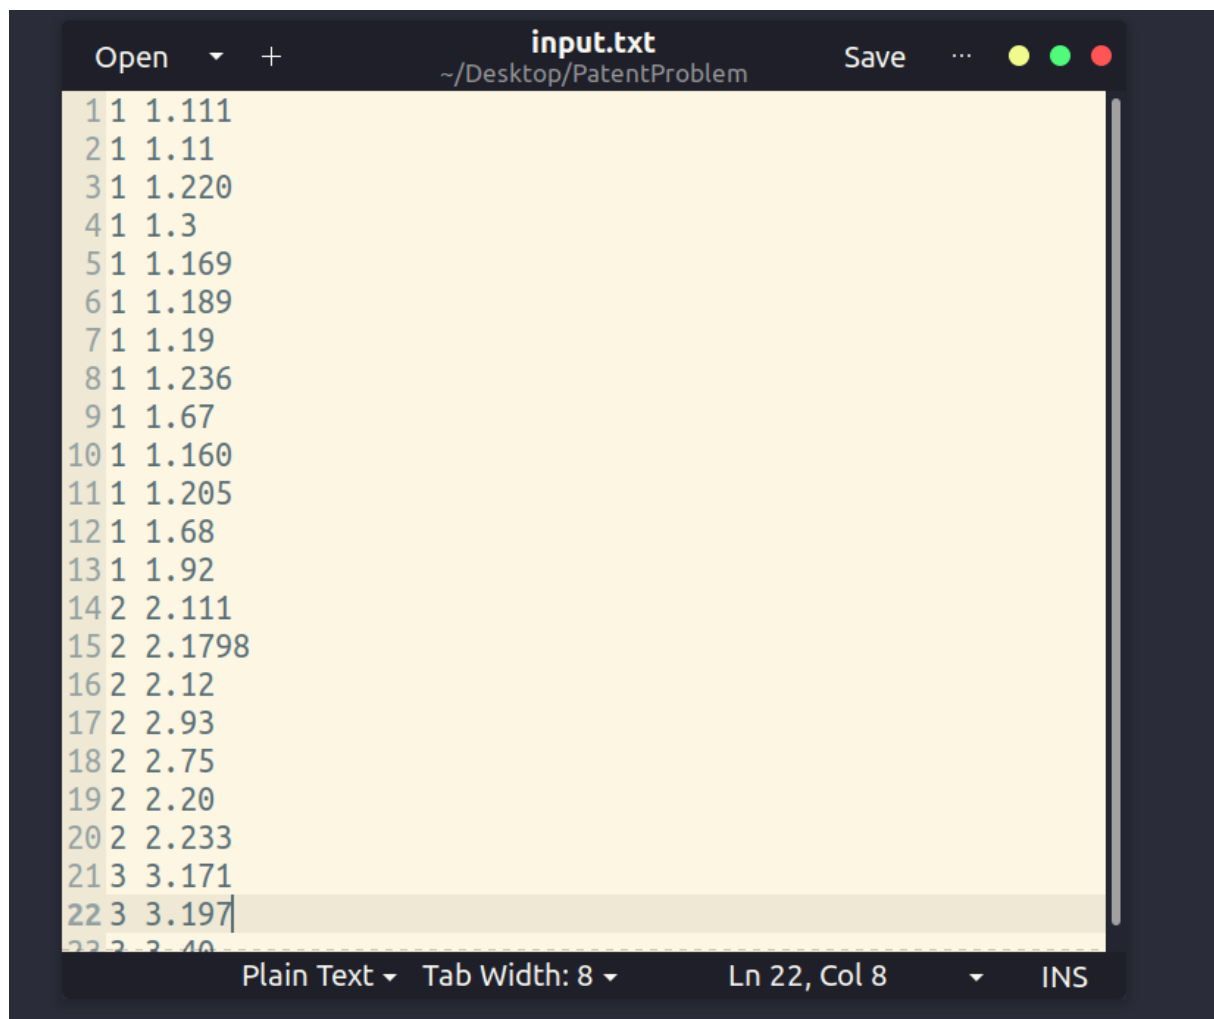
### 1.4.3 Step 3: Create directory structure for program in Hadoop

- We need to create folder to store input data in HDFS by below command:

```
hadoop fs -mkdir /PatentProgram
hadoop fs -mkdir /PatentProgram/Input
hadoop fs -put "local input file's path" /PatentProgram/Input
```

- Example input:

-



```
Open + input.txt Save ...  
~/Desktop/PatentProblem  
1 1 1.111  
2 1 1.11  
3 1 1.220  
4 1 1.3  
5 1 1.169  
6 1 1.189  
7 1 1.19  
8 1 1.236  
9 1 1.67  
10 1 1.160  
11 1 1.205  
12 1 1.68  
13 1 1.92  
14 2 2.111  
15 2 2.1798  
16 2 2.12  
17 2 2.93  
18 2 2.75  
19 2 2.20  
20 2 2.233  
21 3 3.171  
22 3 3.197  
23 3 3.40
```

Plain Text ▾ Tab Width: 8 ▾ Ln 22, Col 8 ▾ INS

**Figure 1.7:** Input file

### 1.4.4 Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" PatentProgram /PatentProgram/Input
↩ /PatentProgram/Output
```

### 1.4.5 Step 5: Final result

- After succesfully calculating, we can check our result in HDFS like below:

The screenshot shows the Hadoop web interface for browsing the HDFS directory. The top navigation bar includes links for Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The main heading is "Browse Directory". Below this, there is a search bar with the path "/PatentProgram/Output" and a "Go!" button. To the right of the search bar are icons for file operations. Below the search bar, there is a "Show" dropdown set to "25" and a "Search:" input field. The main content area displays a table of files and directories. The table has columns for Permission, Owner, Group, Size, Last Modified, Replication, Block Size, Name, and a trash icon. There are two entries in the table: a directory named "\_SUCCESS" and a file named "part-r-00000". Both entries show a permission of "-rw-r--r--", owner "antrvan", group "supergroup", and a replication factor of 1. The file "part-r-00000" has a size of 13 B. Below the table, it says "Showing 1 to 2 of 2 entries". At the bottom right, there are "Previous", "1", and "Next" buttons. The footer of the interface says "Hadoop, 2022."

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	antrvan	supergroup	0 B	Mar 20 15:00	1	128 MB	_SUCCESS
-rw-r--r--	antrvan	supergroup	13 B	Mar 20 15:00	1	128 MB	part-r-00000

Figure 1.8: Output 1

- 
- 

## 1.5 5. MaxTemp Program

### 1.5.1 Step 1: Program's solution

- Import:

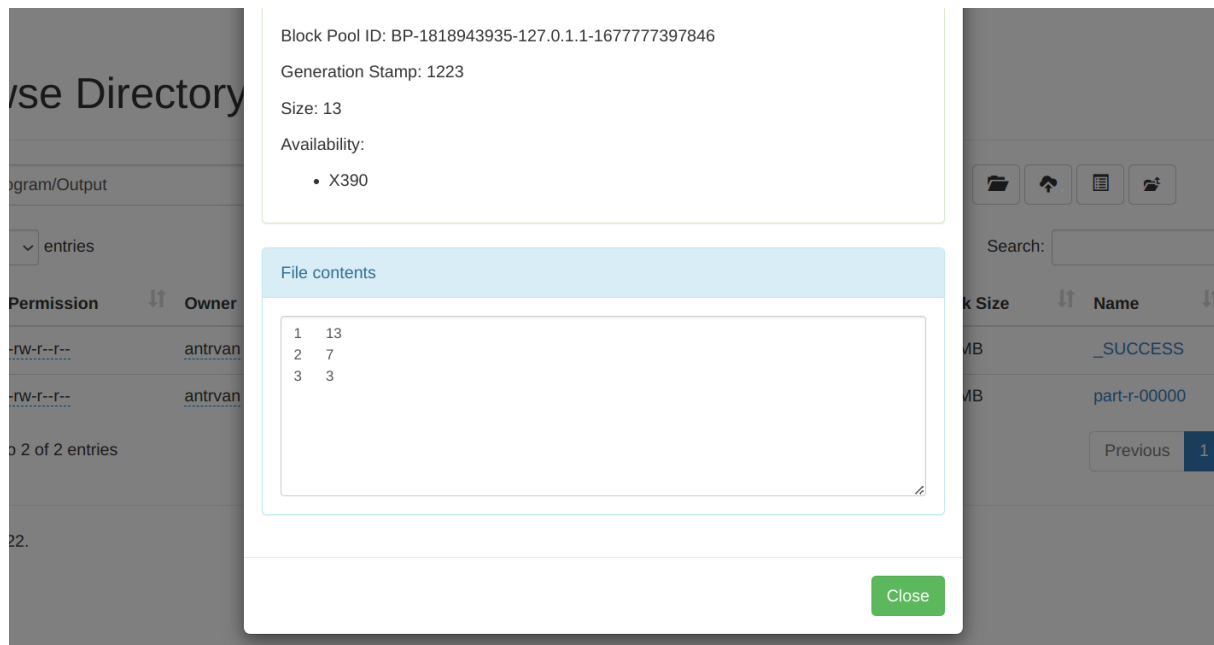


Figure 1.9: Output 2

```
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

- Mapper:

```
public static class MaxTemperatureMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private static final int MISSING = 9999;

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        String line = value.toString();
        String year = line.substring(0, 4);
        int airTemperature = Integer.parseInt(line.substring(5));
        context.write(new Text(year), new IntWritable(airTemperature));
    }
}
```

- Reducer:

```
public static class MaxTemperatureReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int maxTemperature = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            maxTemperature = Math.max(maxTemperature, value.get());
        }
        context.write(key, new IntWritable(maxTemperature));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "max temperature");
    job.setJarByClass(MaxTemp.class);
    job.setMapperClass(MaxTemperatureMapper.class);
    job.setCombinerClass(MaxTemperatureReducer.class);
    job.setReducerClass(MaxTemperatureReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- Explain:
- In the mapper method, we extract the year and temperature from each input line and write them to the key/value pair. We do not need to verify the format of the input stream because in this case all the lines have the same format and we can simply use fixed indexes to extract the information.
- In the reduce method, we find the highest temperature for each year by traversing the list of pooled values for the same key.


### 1.5.2 Step 2: Class Creation

```
jar -cvf MaxTemp.jar -C classes/ .
```

### 1.5.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /MaxTemp
hadoop fs -mkdir /MaxTemp/Input
hadoop fs -put 'local input file's path ' /MaxTemp/Input
```

- Example input:



### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" MaxTemp /MaxTemp/Input /MaxTemp/Output
```

### 1.5.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:
- 
-

/MaxTemp/Output Go! 📁 🔍 📄 📧

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	pmx	supergroup	0 B	Mar 28 22:12	1	128 MB	_SUCCESS	🗑️
<input type="checkbox"/>	-rw-r--r--	pmx	supergroup	24 B	Mar 28 22:12	1	128 MB	part-r-00000	🗑️

Showing 1 to 2 of 2 entries Previous 1 Next

Hadoop, 2022.

Figure 1.10: Output 1

## 1.6 6. AverageSalary Program

### 1.6.1 Step 1: Program's solution

- Mapper:

```
public static class AvgMapper
    extends Mapper<Object, Text, Text, FloatWritable> {

    private Text id = new Text();
    private FloatWritable salary = new FloatWritable();

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String[] values = value.toString().split("\\t");
        id.set(values[0]);

        salary.set(Float.parseFloat(values[2]));
        context.write(id, salary);
    }
}
```

- Reducer

```
public static class AvgReducer
    extends Reducer<Text, FloatWritable, Text, FloatWritable> {

    private FloatWritable result = new FloatWritable();

    public void reduce(Text key, Iterable<FloatWritable> values,
        Context context) throws IOException, InterruptedException {
        float totalSalary = 0;
```

Block information -- Block 0 ▾

Block ID: 1073741991

Block Pool ID: BP-1926325839-127.0.1.1-1678203211420

Generation Stamp: 1167

Size: 24

Availability:

- pmx

File contents

1900	36
1901	48
1902	49

**Figure 1.11:** Output 2



```
int numberPersons = 0;
for (FloatWritable salary : values) {
    totalSalary += salary.get();
    numberPersons++;
}

result.set(totalSalary/numberPersons);
context.write(key, result);
}
```

- Main

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "average salary");
    job.setJarByClass(AverageSalary.class);

    job.setMapperClass(AvgMapper.class);
    job.setReducerClass(AvgReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(FloatWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- In Map function, we will collect employee's ID with their salary to make a pair. Then in Reducer, we will take average salary of each employee's ID to write in the output.

### 1.6.2 Step 2: Class Creation

- After complete code in Java, we need to generate file jar from builded classes by below command:

```
jar -cvf AverageSalary.jar -C classes/ .
```

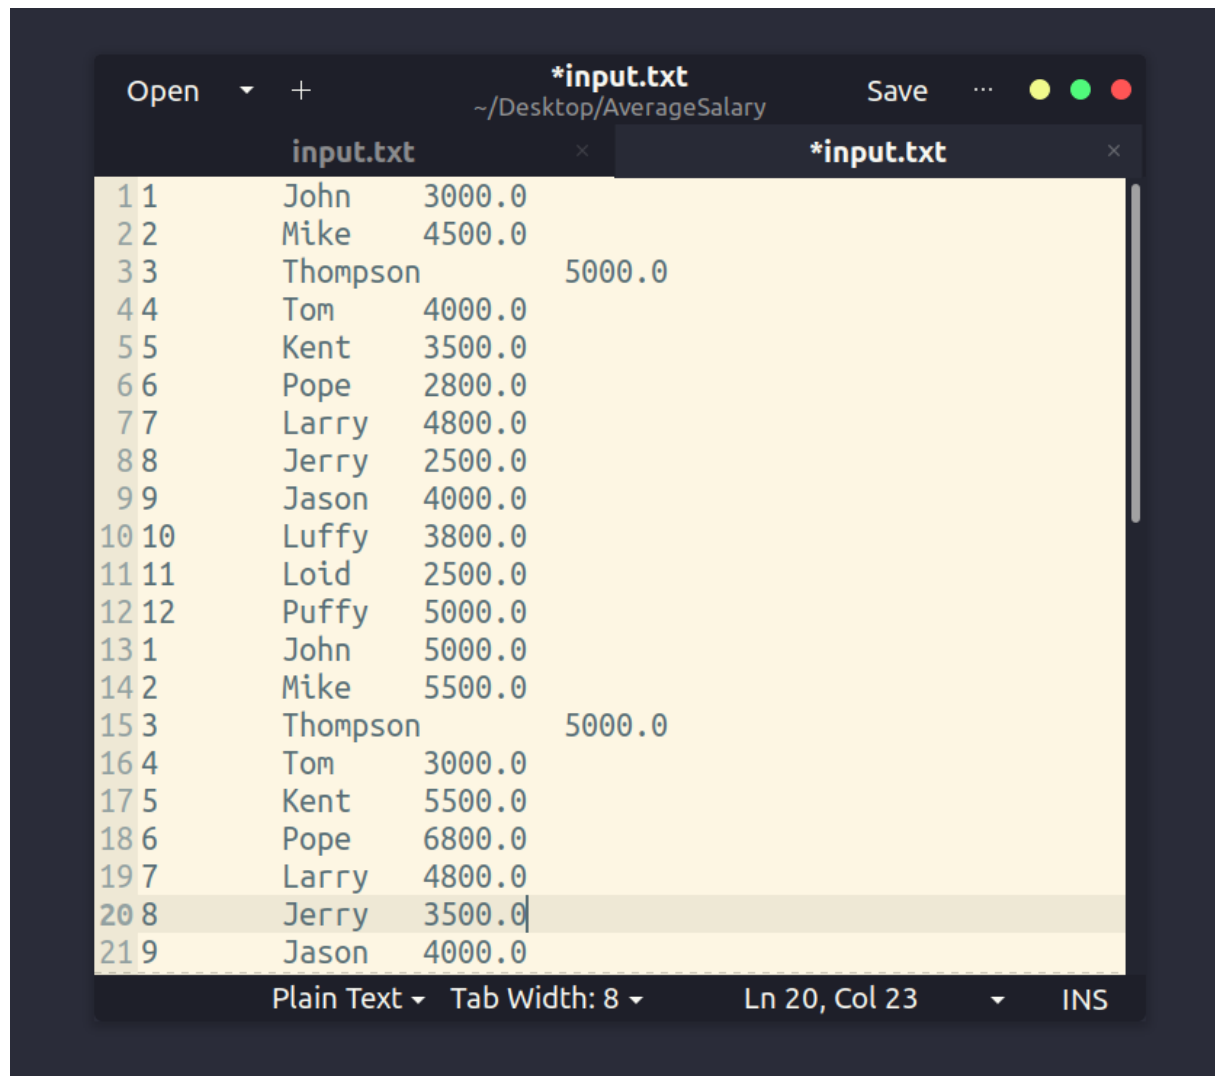
- Notice: Make sure that you export HADOOP\_CLASSPATH before building file jar

### 1.6.3 Step 3: Create directory structure for program in Hadoop

- We need to create folder to store input data in HDFS by below command:

```
hadoop fs -mkdir /AverageSalary
hadoop fs -mkdir /AverageSalary/Input
hadoop fs -put "local input file's path" /AverageSalary/Input
```

- Example input:



**Figure 1.12:** Input file

-

### 1.6.4 Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" AverageSalary /AverageSalary/Input  
↪ /AverageSalary/Output
```

### 1.6.5 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:

The screenshot shows the Hadoop web interface for browsing the directory `/AverageSalary/Output`. The interface includes a navigation bar with links to Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. Below the navigation bar, the title "Browse Directory" is displayed. A search bar and a "Go!" button are present. The main content area shows a table of files and directories. The table has columns for Permission, Owner, Group, Size, Last Modified, Replication, Block Size, Name, and a trash icon. Two entries are listed: `_SUCCESS` (0 B) and `part-r-00000` (111 B). The table is sorted by Name. A "Showing 1 to 2 of 2 entries" message is displayed at the bottom of the table. A pagination bar shows "Previous", "1", and "Next".

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
-rw-r--r--	antivan	supergroup	0 B	Mar 20 15:39	1	128 MB	<a href="#">_SUCCESS</a>	
-rw-r--r--	antivan	supergroup	111 B	Mar 20 15:39	1	128 MB	<a href="#">part-r-00000</a>	

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2022.

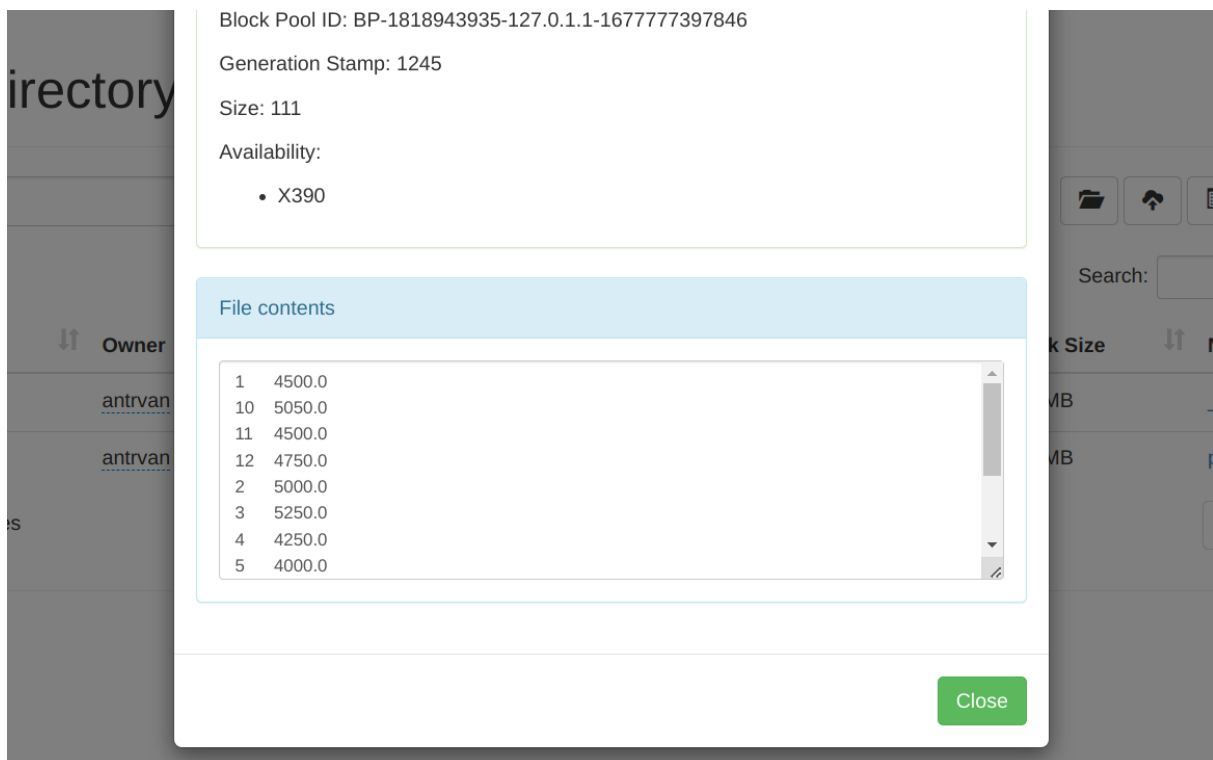
Figure 1.13: Output 1

- 
- 

## 1.7 7. De Identify HealthCare Program

### 1.7.1 Step 3: Program's solution

- Mapper

**Figure 1.14:** Output 2

```
public static Integer[] encryptCol = {2, 3, 4, 5, 6, 7, 8};
private static byte[] key1 = new String("sampleKey1234567").getBytes();

public static class Map
    extends Mapper<Object, Text, NullWritable, Text> {

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        StringTokenizer tokenizer = new StringTokenizer(value.toString(), ",");
        List<Integer> list = new ArrayList<>();

        Collections.addAll(list, encryptCol);
        // list = {2, 3, 4, 5, 6, 7, 8}

        System.out.println("Mapper :: one" + value);
        String newStr = "";

        int counter = 1;

        while (tokenizer.hasMoreTokens()) {
            String token = tokenizer.nextToken();
            System.out.println("token" + token);
            System.out.println("i=" + counter);

            if (list.contains(counter)) {
                if (newStr.length() > 0) {
                    newStr += ",";
                }
                newStr += encrypt(token, key1);
            }
            else {
                if (newStr.length() > 0) {
                    newStr += ",";
                }
                newStr += token;
            }
            counter += 1;
        }

        context.write(NullWritable.get(), new Text(newStr.toString()));
    }
}
```

- Encrypt function

```
public static String encrypt(String strToEncrypt, byte[] key)
{
    try
    {
        Cipher cipher = Cipher.getInstance("AES/ECB/PKCS5Padding");
        SecretKeySpec secretKey = new SecretKeySpec(key, "AES");
        cipher.init(Cipher.ENCRYPT_MODE, secretKey);
    }
}
```

```
String encryptedString =  
    ↪ Base64.encodeBase64String(cipher.doFinal(strToEncrypt.getBytes()));  
  
    return encryptedString.trim();  
}  
catch (Exception e)  
{  
    logger.error("Error while encrypting", e);  
}  
return null;  
}
```

- Main

```
public static void main(String[] args) throws Exception {  
    if (args.length != 2) {  
        System.out.println("usage: [input] [output]");  
        System.exit(-1);  
    }  
  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "de identify data");  
    job.setMapperClass(Map.class);  
  
    job.setInputFormatClass(TextInputFormat.class);  
    job.setOutputFormatClass(TextOutputFormat.class);  
  
    job.setOutputKeyClass(NullWritable.class);  
    job.setOutputValueClass(Text.class);  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.setJarByClass(DeIdentifyData.class);  
    job.waitForCompletion(true);  
}
```

- The idea to resolve this question is only using Map function and encrypt function to encrypt data in identified columns which need to be hidden.

### 1.7.2 Step 2: Class Creation

- After complete code in Java, we need to generate file jar from builded classes by below command:

```
jar -cvf DeIdentifyData.jar -C classes/ .
```

- Notice: Make sure that you export HADOOP\_CLASSPATH before buiding file jar

### 1.7.3 Step 3: Create directory structure for program in Hadoop

- We need to create folder to store input data in HDFS by below command:

```
hadoop fs -mkdir /DeIdentifyData
hadoop fs -mkdir /DeIdentifyData/Input
hadoop fs -put "local input file's path" /DeIdentifyData/Input
```

- Example input:



Figure 1.15: Input file

- 

### 1.7.4 Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" DeIdentifyData /DeIdentifyData/Input
↪ /DeIdentifyData/Output
```

### 1.7.5 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:

Hadoop, 2022.

**Figure 1.16:** Output 1

- 
- 

## 1.8 8 Music Track Program

### 1.8.1 Step 1: Program's solution

- task1: Number of unique listeners
- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobClient;
```



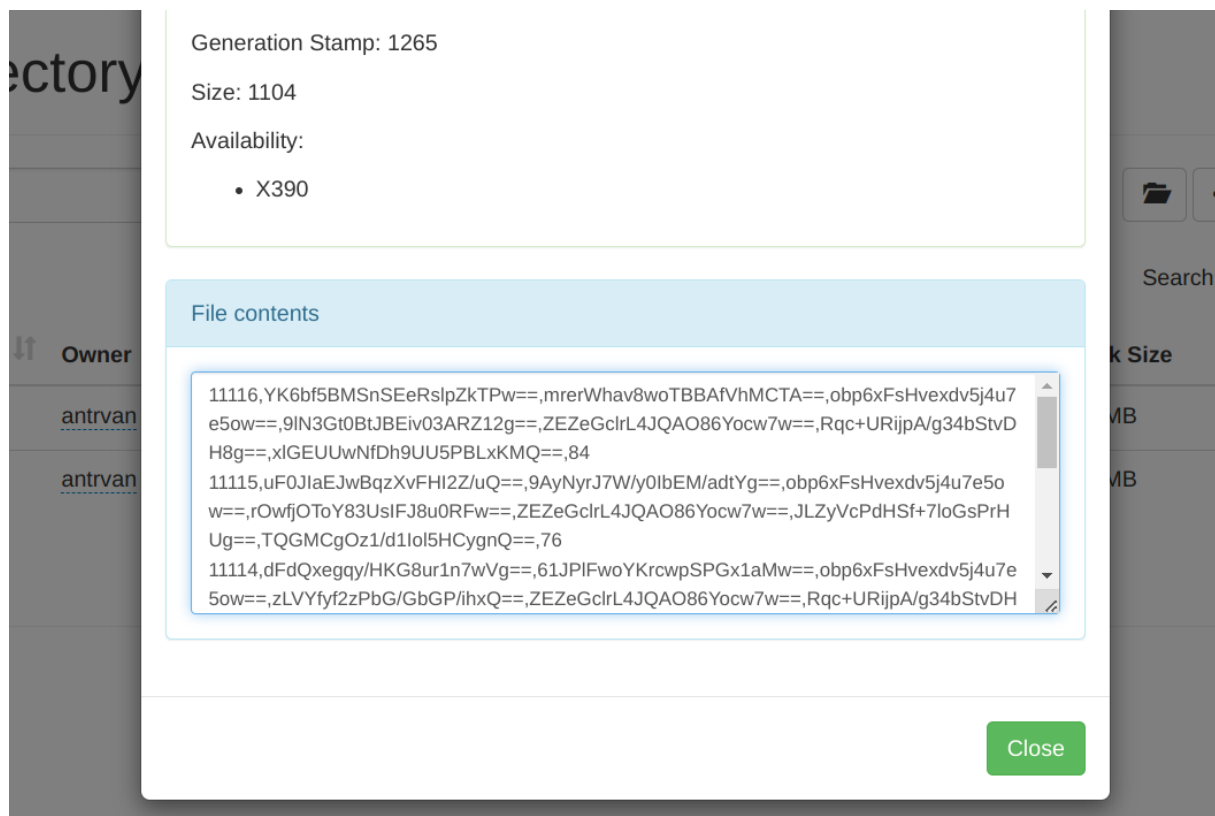


Figure 1.17: Output 2

```
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<Object, Text, IntWritable, IntWritable> {
    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        String[] data = line.split(",");
        IntWritable user = new IntWritable(Integer.parseInt(data[UserId]));
        IntWritable track = new IntWritable(Integer.parseInt(data[TrackId]));
        context.write(track, user);
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<IntWritable, IntWritable, IntWritable,
    ↪ IntWritable> {
    public void reduce(IntWritable key, Iterable<IntWritable> values, Context context)
        ↪ throws IOException, InterruptedException {
        Set<Integer> users = new HashSet<Integer>();
        for (IntWritable val : values) {
            users.add(val.get());
        }
        IntWritable result = new IntWritable(users.size());
        context.write(key, result);
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "Listener");
    job.setJarByClass(Listener.class);
    job.setOutputKeyClass(IntWritable.class);
    job.setOutputValueClass(IntWritable.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
FileOutputFormat.setOutputPath(job, new Path(args[1]));
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
job.waitForCompletion(true);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the String[] to separate the data in the value. Then it sends each pair<trackId,userId> to the Reducer.
- In the reducer method: This reducer will take the pair<trackId,userId> from the Mapper and add it into Set(HashSet) then return <key,Set.size()).
- task2: Number of times the track was shared with others
- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<Object, Text, Text, IntWritable> {
    private Text track = new Text();

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        String[] data = line.split(",");
        track = new Text(data[TrackId]);
        context.write(track, new IntWritable(Integer.valueOf(data[Shared])));
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum = sum + val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "Shared");
    job.setJarByClass(Shared.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);
    job.waitForCompletion(true);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the String[] to separate the data in the value. Then it sends each pair<trackId,Shared> to the Reducer.
- In the reducer method: This reducer will take the pair<trackId,Shared> from the Mapper and calculate the total number of each track was shared with others by adding the values of 1s together..
- task3: Number of times the track was listened to on the radio
- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<Object, Text, Text, IntWritable> {
    private Text track = new Text();

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        String[] data = line.split(",");
        track = new Text(data[TrackId]);
        context.write(track, new IntWritable(Integer.valueOf(data[Radio])));
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum = sum + val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "Radio");
    job.setJarByClass(Radio.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    job.setMapperClass(Map.class);
}
```

```
job.setCombinerClass(Reduce.class);
job.setReducerClass(Reduce.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
job.waitForCompletion(true);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the String[] to separate the data in the value. Then it sends each pair<trackId, Radio> to the Reducer.
- In the reducer method: This reducer will take the pair<trackId, Radio> from the Mapper and calculate the total number of each track was listened to on radio by adding the values of 1s together..
- task4: Number of times the track was listened to in total
- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<Object, Text, Text, IntWritable> {
    private Text track = new Text();

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        String[] data = line.split(",");
```

```
        track = new Text(data[TrackId]);
        context.write(track, new IntWritable(Integer.valueOf(data[Skip])));
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            if (val.get() == 0)
                sum++;
        }
        context.write(key, new IntWritable(sum));
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "Total");
    job.setJarByClass(Total.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);
    job.waitForCompletion(true);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the String[] to separate the data in the value. Then it sends each pair<trackId,Skip> to the Reducer.
- In the reducer method: This reducer will take the pair<trackId,Skip> from the Mapper and calculate the total number of each track was'n skipped by adding the values of 1s together..
- task5: Number of times the track was skipped on the radio
- Import:

```
import java.util.*;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

- Mapper:

```
public static class Map extends Mapper<Object, Text, Text, IntWritable> {
    private Text track = new Text();

    public void map(Object key, Text value, Context context) throws IOException,
        ↪ InterruptedException {
        String line = value.toString();
        String[] data = line.split(",");
        track = new Text(data[TrackId]);
        if (data[Radio].equals("0") || data[Skip].equals("0"))
            context.write(track, new IntWritable(0));
        else
            context.write(track, new IntWritable(1));
    }
}
```

- Reducer:

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum = sum + val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

- Main:



```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "Skip_Radio");
    job.setJarByClass(Skip_Radio.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);
    job.waitForCompletion(true);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the String[] to separate the data in the value. Then it sends each pair<trackId,Skip&Radio> to the Reducer.
- In the reducer method: This reducer will take the pair<trackId,Skip&Radio> from the Mapper and calculate the total number of each track was skipped on the radio by adding the values of 1s together..


### 1.8.2 Step 2: Class Creation

```
jar -cvf fileName.jar -C classes/ .
```

### 1.8.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /fileName
hadoop fs -mkdir /fileName/Input
hadoop fs -put 'local input file's path ' /fileName/Input
```

- Example input:







### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" WordCount /WordCount/Input /WordCount/Output
```



#### 1.8.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:
- task1:
-

## Browse Directory

/musictrack Go!    

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	thien123456	supergroup	0 B	Mar 31 17:39	0	0 B	input	
<input type="checkbox"/>	drwxr-xr-x	thien123456	supergroup	0 B	Mar 31 17:52	0	0 B	output	

Showing 1 to 2 of 2 entries Previous 1 Next

Hadoop, 2022.

Figure 1.18: Output 1

Hadoop Overview Datanodes DataNode Volume Failures Snapshot Startup Progress Utilities

## Browse Directory

/musictrack/output

Show 25 entries

<input type="checkbox"/>	Permission	Owner
<input type="checkbox"/>	-rw-r--r--	thien123
<input type="checkbox"/>	-rw-r--r--	thien123

Showing 1 to 2 of 2 entries

Hadoop, 2022.

File information - part-r-00000

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073742078  
Block Pool ID: BP-1163893533-127.0.1.1-1679742735770  
Generation Stamp: 1254  
Size: 18  
Availability:  
• thien

File contents

```
222 1
223 1
225 2
```

Close

- 
- task2:
-

File information - part-r-00000

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073742017

Block Pool ID: BP-1163893533-127.0.1.1-1679742735770

Generation Stamp: 1193

Size: 18

Availability:

- thien

File contents

```
222 1
223 1
225 1
```

Close

**Figure 1.19:** Output 1

The screenshot shows a Hadoop web interface with a modal window open. At the top of the modal, there are three tabs: "Download", "Head the file (first 32K)", and "Tail the file (last 32K)". Below these tabs is a green header bar with the text "Block information --" and a dropdown menu currently set to "Block 0". The main content area of the modal displays the following information:

- Block ID: 1073742038
- Block Pool ID: BP-1163893533-127.0.1.1-1679742735770
- Generation Stamp: 1214
- Size: 18
- Availability:
  - thien

Below this information is a section titled "File contents" with a light blue header. It contains a text area displaying the following lines:

```
222 0
223 0
225 2
```

At the bottom right of the modal is a green "Close" button. The background of the interface is dimmed, showing a file browser with columns for "Owner", "Block Size", and "Name".

- 
- task3:
-

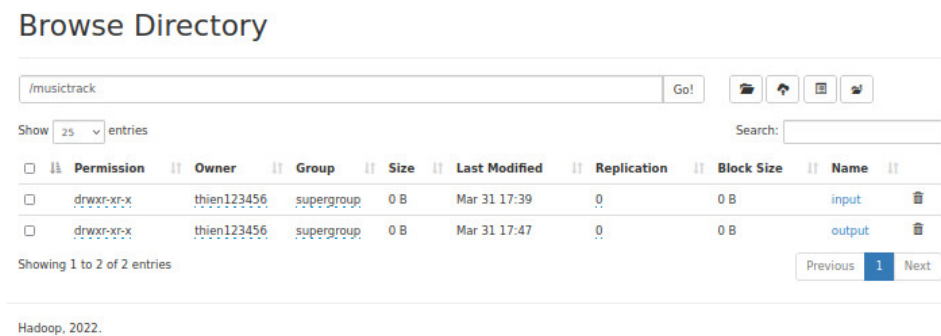
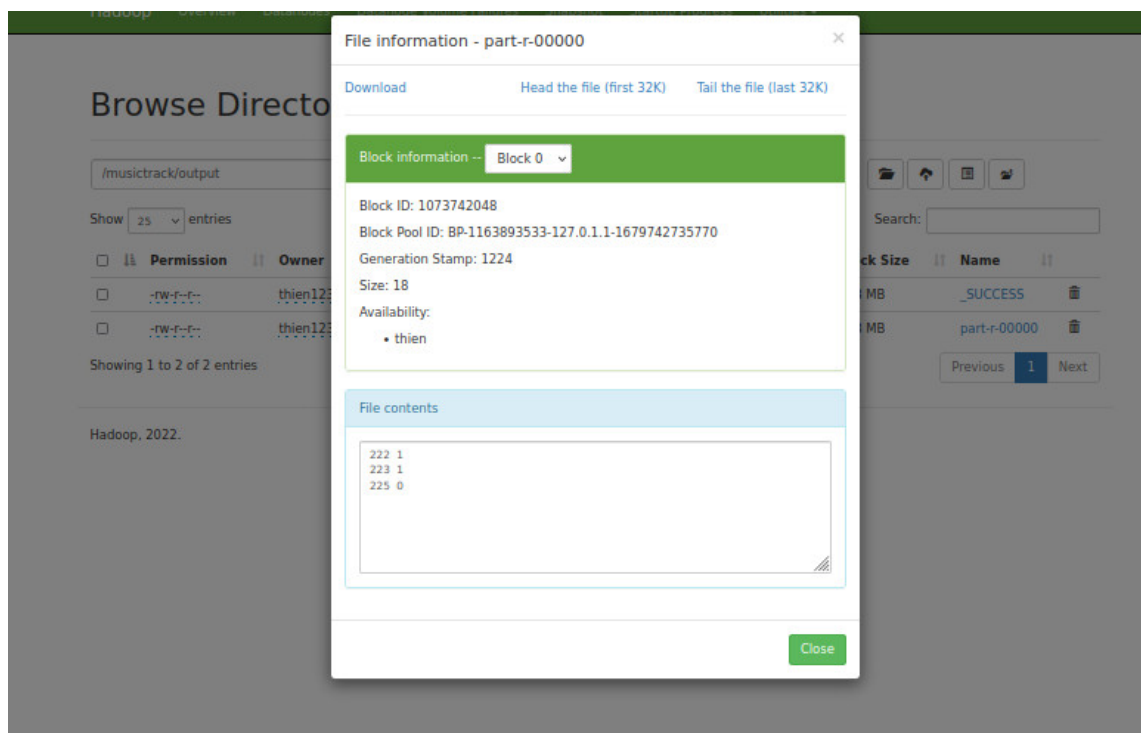






Figure 1.20: Output 1





- 
- task4:
-

## Browse Directory

/musictrack/output Go!    

Show  entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	thien123456	supergroup	0 B	Mar 31 17:49	1	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	thien123456	supergroup	18 B	Mar 31 17:49	1	128 MB	part-r-00000	

Showing 1 to 2 of 2 entries Previous **1** Next

Hadoop, 2022.

Figure 1.21: Output 1

File information - part-r-00000

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073742058  
Block Pool ID: BP-1163893533-127.0.1.1-1679742735770  
Generation Stamp: 1234  
Size: 18  
Availability:  
• thien

File contents

```
222 0
223 1
225 0
```

Close

- 
- task5:
- 
-

## Browse Directory

/musictrack

Show  entries

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	thien123456	supergroup	0 B	Mar 31 17:39	0	0 B	input	
<input type="checkbox"/>	drwxr-xr-x	thien123456	supergroup	0 B	Mar 31 17:51	0	0 B	output	

Showing 1 to 2 of 2 entries

Hadoop, 2022.

Figure 1.22: Output 1

Browse Directory

/musictrack/output

Show  entries

<input type="checkbox"/>	Permission	Owner
<input type="checkbox"/>	-rw-r--r--	thien123
<input type="checkbox"/>	-rw-r--r--	thien123

Showing 1 to 2 of 2 entries

Hadoop, 2022.

File information - part-r-00000

Block information --

Block ID: 1073742068  
Block Pool ID: BP-1163893533-127.0.1.1-1679742735770  
Generation Stamp: 1244  
Size: 18  
Availability:  
• thien

File contents

```
222 0
223 1
225 0
```

Figure 1.23: Output 2



## 1.9 9. Telecom Call Data Record Program

### 1.9.1 Step 1: Program's solution

- Import:

```
import java.io.IOException;
import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.Date;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

- Mapper:

```
public static class Map
    extends Mapper<Object, Text, Text, LongWritable>{

    Text phoneNumber = new Text();
    LongWritable minutes = new LongWritable();

    public void map(Object key, Text value, Context context
    ) throws IOException, InterruptedException {
        String line = value.toString();
        String[] tokens = line.split("\\|");
        if(tokens[4].equals("1")) {
            phoneNumber.set(tokens[0]);
            minutes.set(calculateTimeInMinutes(tokens[2], tokens[3]));
            context.write(phoneNumber, minutes);
        }
    }

    private long calculateTimeInMinutes(String start, String end) {
        SimpleDateFormat formatter = new SimpleDateFormat(("yyyy-MM-dd HH:mm:ss"));
        long minutes = -1; // if this value happen then there's an error
        try {
            // put code in try catch so that java is not angry
            Date startDate = formatter.parse(start);
            Date endDate = formatter.parse(end);
```

```
        long duration = endDate.getTime() - startDate.getTime();
        minutes = duration / (1000 * 60);
    } catch (ParseException e) {
        e.printStackTrace();
    }
    return minutes;
}
}
```

- Reducer:

```
public static class Reduce
    extends Reducer<Text,LongWritable,Text,LongWritable>{

    public void reduce(Text key, Iterable<LongWritable> values,
        Context context
    ) throws IOException, InterruptedException {
        long totalMinutes = 0;
        for(LongWritable val: values) {
            totalMinutes += val.get();
        }
        if(totalMinutes > 60) {
            context.write(key, new LongWritable(totalMinutes));
        }
    }
}
```

- Main:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "call data record");
    job.setJarByClass(CallDataRecord.class);
    job.setMapperClass(Map.class);
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(LongWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

- Explain:
- In the mapper method: This mapper will take as input an Object (representing the key), and a Text (representing the value) and use the `split()` method to separate the fields in a line (inside is a regular expression indicate a “|” character). It then check if the std is equal to 1. If true, it will write

the FromPhoneNumber and the time in minute (calculate using the calculateTimeInMinutes() utility function) to the context.

- In the reducer metho: This reducer will take the words from the Mapper and add all the values together, it then write the result to the context if the total is greater than 60.

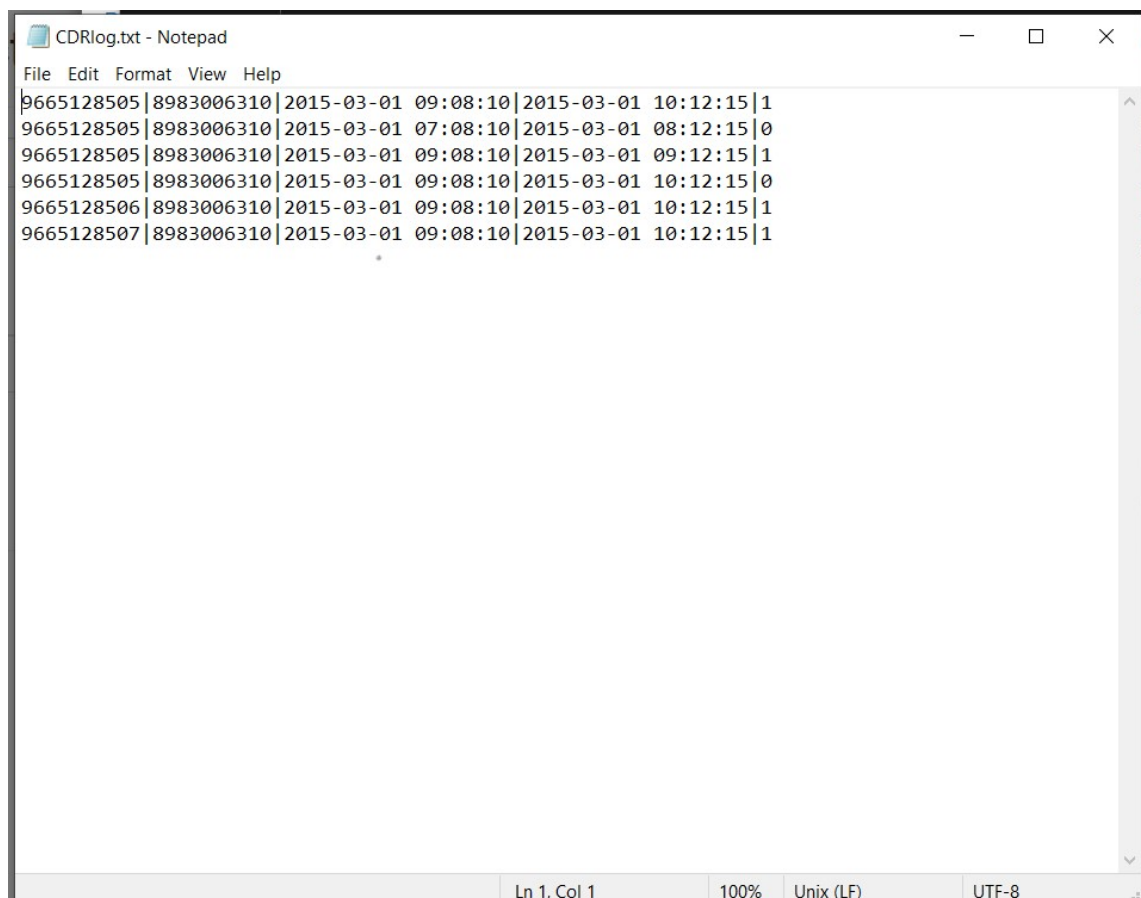
### 1.9.2 Step 2: Class Creation

```
jar -cvf CallDataRecord.jar -C classes/ .
```

### 1.9.3 Step 3: Create directory structure for program in Hadoop

```
hadoop fs -mkdir /Phone
hadoop fs -mkdir /Phone/Input
hadoop fs -put 'local input file's path ' /Phone/Input
```

- Example input:



### ### Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" CallDataRecord /Phone/Input /Phone/Output
```

## 1.9.4 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:

### Browse Directory

/Phone/Output

Go!

Show

25

entries

Search:

<input type="checkbox"/>		Permission		Owner		Group		Size		Last Modified		Replication		Block Size		Name	
<input type="checkbox"/>		-rw-r--r--		ADMIN		supergroup		0 B		Mar 31 21:12		1		128 MB		<a href="#">_SUCCESS</a>	
<input type="checkbox"/>		-rw-r--r--		ADMIN		supergroup		42 B		Mar 31 21:12		1		128 MB		<a href="#">part-r-00000</a>	

Showing 1 to 2 of 2 entries

Previous

1

Next

Hadoop, 2022.

Hadoop, 2022.

**Figure 1.24:** Output 1

- 
- 

## 1.10 10. Count Connected Components Program

### 1.10.1 Step 1: Program's solution

- Mapper

```
public static class Map
    extends Mapper<Object, Text, Text, Text> {

    public void map(Object key, Text value, Context context
    ) throws IOException, InterruptedException {
        String[] tokens = value.toString().split(" ");
```

The screenshot shows a web interface for viewing file information. At the top, there's a title bar 'File information - part-r-00000' with a close button. Below the title bar, there are three tabs: 'Download', 'Head the file (first 32K)', and 'Tail the file (last 32K)'. The 'Download' tab is active. Underneath, there's a green header bar with 'Block information --' and a dropdown menu showing 'Block 0'. Below this, the block details are listed: 'Block ID: 1073742130', 'Block Pool ID: BP-938757607-192.168.1.3-1679730241008', 'Generation Stamp: 1306', 'Size: 42', and 'Availability:'. Under 'Availability', there's a list with one item: 'LAPTOP-1M9L4SKS'. Below the block information, there's a light blue header bar with 'File contents'. Underneath, there's a text area showing three lines of data: '9665128505 68', '9665128506 64', and '9665128507 64'. A small icon is visible in the bottom right corner of the text area.

File information - part-r-00000

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073742130  
Block Pool ID: BP-938757607-192.168.1.3-1679730241008  
Generation Stamp: 1306  
Size: 42  
Availability:

- LAPTOP-1M9L4SKS

File contents

```
9665128505 68
9665128506 64
9665128507 64
```

**Figure 1.25:** Output 2

```
String keyValue = tokens[0];
Arrays.sort(tokens);

int i = 0;
while (i < tokens.length) {
    context.write(new Text("map"), new Text(keyValue + "," + tokens[i]));
    i++;
}
}
```

- Reducer

```
public static class Reduce
    extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values,
        Context context
    ) throws IOException, InterruptedException {

        TreeMap<Integer, ArrayList<Integer>> sortedMap = new TreeMap<>();
        HashMap<Integer, Integer> result = new HashMap<>();
        for (Text value : values) {
            String[] pair = value.toString().split(",");

            int keyItem = Integer.parseInt(pair[0]);
            int valueItem = Integer.parseInt(pair[1]);

            ArrayList<Integer> tmp = sortedMap.getOrDefault(keyItem, new
                ArrayList<Integer>());
            tmp.add(valueItem);
            Collections.sort(tmp);
            sortedMap.put(keyItem, tmp);
        }

        for (Integer k : sortedMap.keySet()) {
            Integer start = sortedMap.get(k).get(0);
            if (start.compareTo(k) == 0) {
                result.put(k, k);
            }
            if (start.compareTo(k) < 0) {
                result.put(k, result.get(start));
                for (Integer v : sortedMap.get(k)) {
                    if (v.equals(start)) continue;
                    for (Integer j : result.keySet()) {
                        if (result.get(j).equals(v)) {
                            result.replace(j, v, start);
                        }
                    }
                }
            }
        }
    }
}
```

```

    }

    HashSet<Integer> components = new HashSet<>();
    boolean b = components.addAll(result.values());

    if (b) {
        context.write(new Text(""), new Text(String.valueOf(components.size())));
    }
}
}

```

- Main

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "count connected component program");

    job.setJarByClass(CountConnectedComponentProgram.class);

    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}

```

- This task is an intriguing question that calculate numbers of separated components in a graph. To resolve this problem, we put pair of source and destination point of every edges in graph to reducer. We put all pairs to TreeMap to sort them. Then in each components, we mark all connected vertices value to smallest vertex. Finally, the result equals numbers of different values in HashMap.

### 1.10.2 Step 2: Class Creation

- After complete code in Java, we need to generate file jar from builded classes by below command:

```

jar -cvf CountConnectedComponentProgram.jar -C classes/ .

```

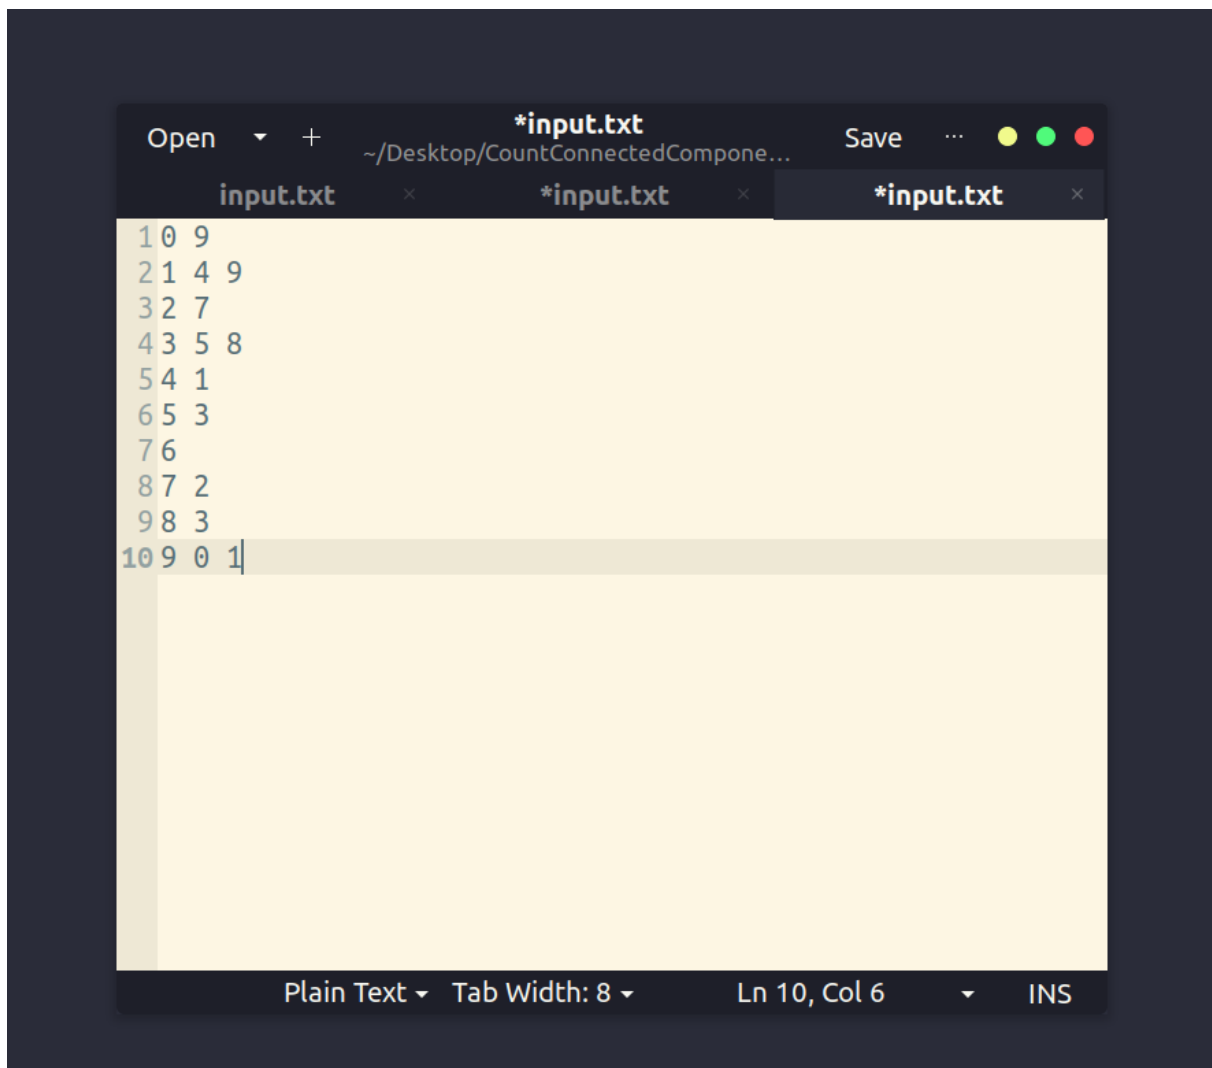
- Notice: Make sure that you export HADOOP\_CLASSPATH before buiding file jar

### 1.10.3 Step 3: Create directory structure for program in Hadoop

- We need to create folder to store input data in HDFS by below command:

```
hadoop fs -mkdir /CountConnectedComponentProgram
hadoop fs -mkdir /CountConnectedComponentProgram/Input
hadoop fs -put "local input file's path" /CountConnectedComponentProgram/Input
```

- Example input:



**Figure 1.26:** Input file

-

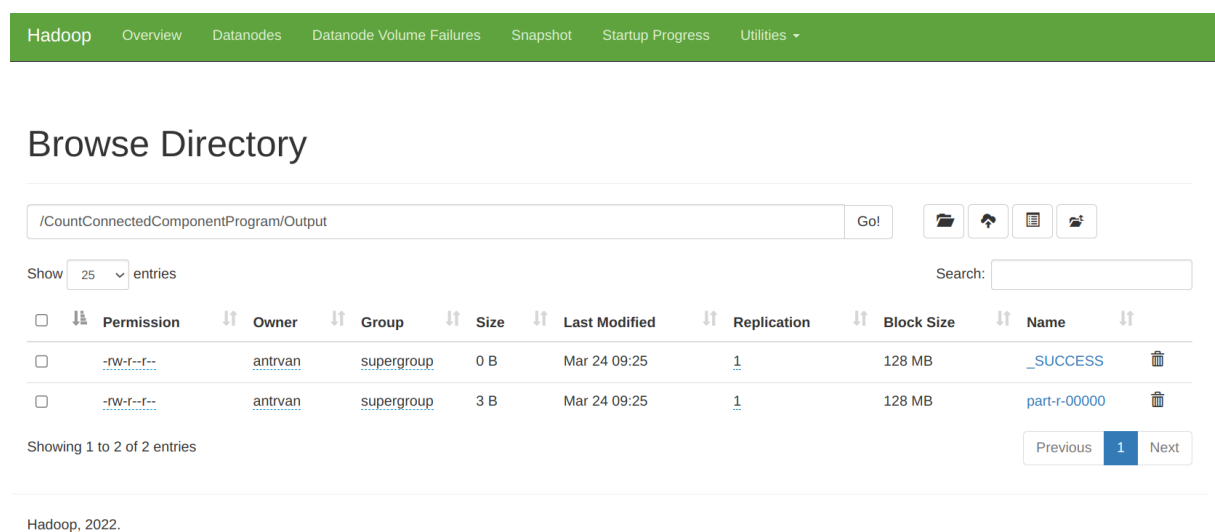


### 1.10.4 Step 4: Create Jar File and deploy it to Hadoop

```
hadoop jar "Path to your local file .jar" CountConnectedComponentProgram
↪ /CountConnectedComponentProgram/Input /CountConnectedComponentProgram/Output
```

### 1.10.5 Step 5: Final result

- After successfully calculating, we can check our result in HDFS like below:



Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/CountConnectedComponentProgram/Output Go!

Show  entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	antrvan	supergroup	0 B	Mar 24 09:25	1	128 MB	<a href="#">_SUCCESS</a>	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	antrvan	supergroup	3 B	Mar 24 09:25	1	128 MB	<a href="#">part-r-00000</a>	<input type="checkbox"/>

Showing 1 to 2 of 2 entries

Previous **1** Next

Hadoop, 2022.

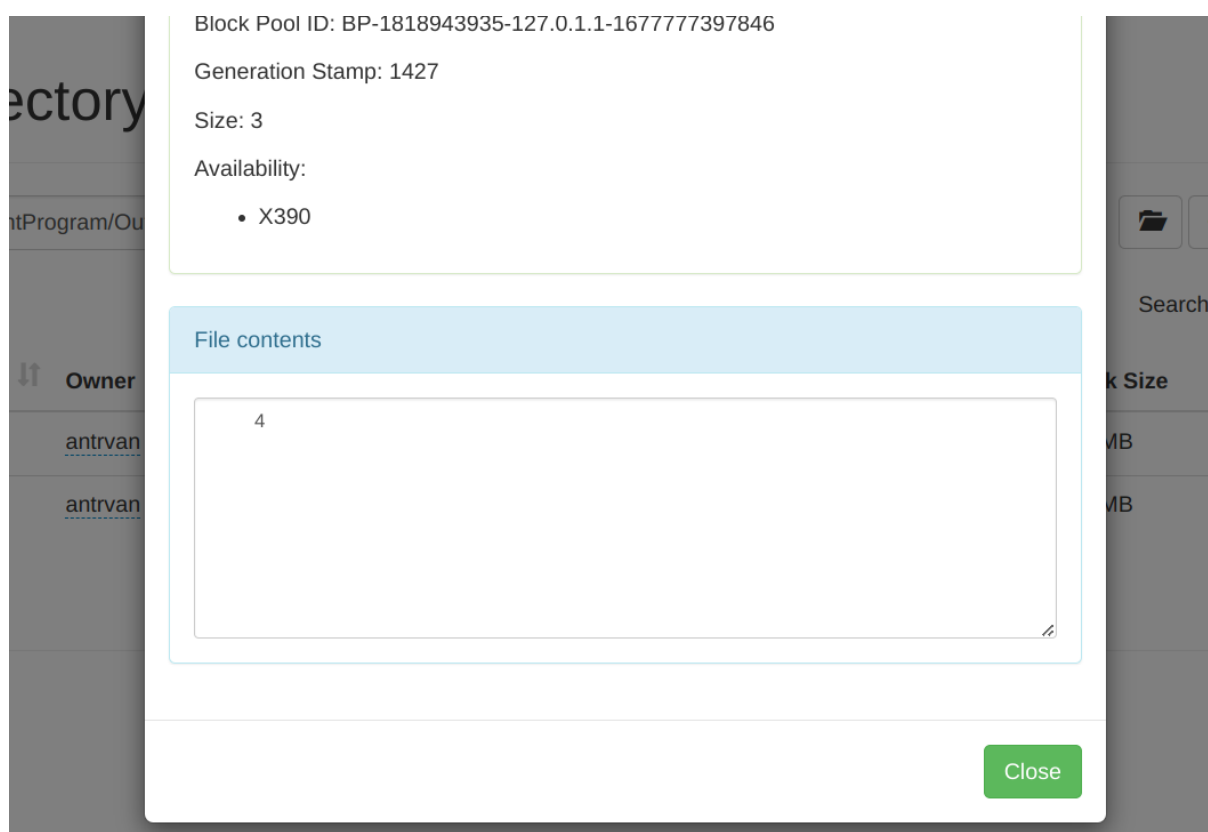
**Figure 1.27:** Output 1

- 
- 

## 1.11 Self-reflection

### 1.11.1 20127435 - Tran Van An

- After completing above tasks, I know more about the useful of MapReduce in real-problems in many aspects as well as get experiences in MapReduce Programing for the midterm test.

**Figure 1.28:** Output 2

**1.11.2 20127395 - Phan Minh Xuan**

- After completing above tasks, I understand how to store, process and manage large data sets, develop skills in the field of big data, especially know more about java language.

**1.11.3 20127032 - Bui Gia Huy**

- After completing above tasks, I know how to set up and manipulate a basic map reduce program, as well as transforming data using java utility class, as well as familiarize myself with java syntax.

**1.11.4 20127631 - Thai Van Thien**

- After completing the above tasks, I know how to set up and work with a basic map reduction program and have a preparation for the midterm exam. ## Member's contribution

Task	Result
1.WordCount Program	100%
2.WordSizeWordCount Program	100%
3.Weather Data	100%
4.PatentProgram	100%
5.MaxTemp Program	100%
6.Average Salary	100%
7.De Identify Data	100%
8.Music Track Program	100%
9.Telecom Call Data Record Program	100%
10.Count Connected Components	100%

MSSV	Member	Contribution Percentage
20127435	Tran Van An	25%
20127395	Phan Minh Xuan	25%
20127032	Bui Gia Huy	25%

MSSV	Member	Contribution Percentage
20127631	Thai Van Thien	25%

## 1.12 References

- Lab requirement pdf